

Sofia Voice Pipeline

How a comment in Sofia Hsu's live room becomes a spoken reply in her own cloned voice — end to end, running on one local GPU.


LIVE at :8080

reply time ~20s

100% local · no cloud calls

voice polish off

WORKFLOW



Five hops from a typed comment to a spoken reply. The dashed hop is a real, trained step in the pipeline — RVC timbre conversion — switched off because on this GPU it roughly triples wall-clock time for a modest gain in voice-identity match.

STACK

Ollama
model runtime

Qwen3-8B
persona LLM

CosyVoice2
zero-shot TTS

RVC
voice polish · off

PyTorch
inference

Tesla V100
16GB · 1 GPU

SUMMARY

Viewers comment in Sofia's live room; a local **Qwen3-8B** model writing through Ollama drafts her reply in character, **CosyVoice2** clones her voice from a 6.5-second reference to speak it, and the room plays the clip back — no data leaves the machine. An **RVC** voice-conversion pass can sharpen the timbre match further but roughly triples wall-clock time on this single GPU, so it ships switched off. End-to-end reply time is **about 20 seconds**, set mostly by CosyVoice2's own synthesis speed rather than any one avoidable delay — reaching sub-3-second replies would need swapping in a streaming-native TTS engine, not further tuning this one.